



One-shot survey marginals: configurations with all 100 valid answers


Same 100 questions; lower is better




Error spread: 10-90%, thick 25-75%, dot median



Paired mean differences and 95% family-bootstrap intervals
Negative favors the one-shot model; intervals are not multiplicity-adjusted




Accuracy and observed cost: complete runs only




#	Model	EP cost
1	GPT-6 Astra (medium)	\$1.588
2	GPT-6 Astra (high)	\$3.287
3	Fable 5.1	\$2.908
4	Claude Opus 5	\$1.058
5	Gemini 3.8 Flash	\$3.456
6	GPT-5.6 Sol	\$1.797
7	Gemini 3.1 Pro preview	\$2.339
8	Grok 4.6	\$2.303
9	GPT-5.6 Luna	\$0.124
10	Claude Sonnet 5	\$0.280
11	Gemini 3.5 Flash Lite	\$0.259
12	MiniMax M3	\$0.104
13	GPT-5.4 mini	\$1.543
14	DeepSeek V4.1 Flash	\$0.017
15	Grok 4.20 reasoning	\$0.547
16	DeepSeek V4 Pro	\$0.114
17	Gemma 4 31B	\$0.012

Valid-answer coverage: every benchmark configuration




GPT-6 reasoning effort: paired mean difference +0.20 TVD (95% CI -0.30, +0.64)



Each dot is the same question at both efforts




Negative favors high; positive favors medium




Survey GPT-6 Astra (high) Fable 5.1
Aaru GPT-6 Astra (medium) Gemini 3.8 Flash




Source: www.philadelphiafed.org




Source: article.images.consumerreports.org



Source: article.images.consumerreports.org



Source: ag.purdue.edu



Source: files.kff.org

Source: api-test.yougov.com